

# Three paradigms for modeling: curve-fitting, probabilistic models and uncertainty quantification, illustrated using the MMI data

By Charles Zheng, using data from Keren et al. (2020)

## Importing the data

```
In [1]: import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
import pandas as pd
```

```
In [2]: # Import the data
ts = pd.read_csv('../mturk.csv', index_col=0)
```

```
In [3]: # Columns
ts.columns
```

```
Out[3]: Index(['CertainAmount', 'Outcome1', 'Outcome2', 'isHappyBlock', 'Happiness',
              'Outcome', 'time', 'subject_id', 'InterpHappiness', 'MoodTarget',
              'Expected', 'Actual', 'CertCol', 'ExpectedCol', 'RPECol', 'Gamble',
              'RutC', 'RutE', 'RutR', 'block'],
              dtype='object')
```

```
In [4]: n_trials = 81
n_subjects = int(len(ts)/n_trials)
# reformat data into separate variables that are #trials x #subjects
all_trial_nos = ts.time.values.reshape((-1, n_trials)).T
all_participant = ts.subject_id.values.reshape((-1, n_trials)).T
all_outcome1 = ts.Outcome1.values.reshape((-1, n_trials)).T
all_outcome2 = ts.Outcome2.values.reshape((-1, n_trials)).T
all_certainAmount = ts.CertainAmount.values.reshape((-1, n_trials)).T
all_choice = ts.Gamble.values.reshape((-1, n_trials)).T
all_outcomeAmount = ts.Actual.values.reshape((-1, n_trials)).T
all_mood_rating = ts['Happiness'].values.reshape((-1, n_trials)).T
n_subjects
```

```
Out[4]: 80
```

```
In [5]: all_winAmount = np.maximum(all_outcome1, all_outcome2)
all_loseAmount = np.minimum(all_outcome1, all_outcome2)
```

# I. Curve-fitting

The LTA model with time drift is defined as follows

$$E(t) = \frac{1}{t-1} \sum_{u=1}^{t-1} A(u)$$

$$\hat{M}(t) = M_0 + \beta_E \sum_{u=1}^t \lambda^{t-u} E(u) + \beta_A \sum_{u=1}^t \lambda^{t-u} A(u) + \beta_T T(t)$$

where  $A(t)$  is the actual outcome of trial  $t$ ,  $T(t)$  is the time stamp (here just the trial number),  $M(t)$  is the mood rating.

$\lambda$  and  $M_0$  are constrained to lie in  $[0,1]$ .  $\beta_A$  and  $\beta_E$  are constrained to be nonnegative. There are no constraints on  $\beta_T$ .

```
In [182]: ## Define a class for the LTA model

class CurveLTA(object):

    def __init__(self):
        pass


    # initializes with some default parameters
    def initialize(self):
        self.m0 = 0.5
        self.lam = 0.8
        self.betaE = 0.01
        self.betaA = 0.005
        self.betaT = 0.0001

    def predict(self, actual, timestamps):
        n_trials = len(actual)
        # holds the predicted moods
        mood_pred = np.zeros(n_trials)
        # Holds the exponentially weighted sums for E(t) and A(t)
        sum_E = 0
        sum_A = 0
        for trial_no in range(n_trials):
            if trial_no == 0:
                lte = 0
            else:
                lte = np.mean(actual[:trial_no])
            sum_E = sum_E * self.lam + lte
            sum_A = sum_A * self.lam + actual[trial_no]
            mood_mu = self.m0 + self.betaE * sum_E + self.betaA * sum_A
            + self.betaT * timestamps[trial_no]
            mood_pred[trial_no] = mood_mu
        return mood_pred
```

```
In [197]: # use a single subject for all demonstrations
subject_index = 7
time = all_trial_nos[:, subject_index]
actual = all_outcomeAmount[:, subject_index]
mood = all_mood_rating[:, subject_index]
highGamble = all_winAmount[:, subject_index]
lowGamble = all_loseAmount[:, subject_index]
certain = all_certainAmount[:, subject_index]
choice = all_choice[:, subject_index]

CL = CurveLTA()
CL.initialize()
plt.plot(time, CL.predict(actual, time), c="r")
plt.scatter(time, mood, c="b")
```

Out[197]: <matplotlib.collections.PathCollection at 0x7fc0e0f51d00>



## Defining a loss function for optimization

```
In [198]: from scipy.optimize import minimize, Bounds
```

```
In [199]: CL = CurveLTA()
def loss_func(par):
    CL.m0, CL.lam, CL.betaE, CL.betaA, CL.betaT = par
    mood_pred = CL.predict(actual, time)
    return np.nansum(np.abs(mood - mood_pred))
```

```
In [200]: # What's the loss on the default parameters?
loss_func([0.5, 0.8, 0.01, 0.005, 0.0001])
```

Out[200]: 6.919534885155892

```
In [201]: # call to minimize the function
minimize(loss_func, [0.5,0.8,0.01,0.005,0.0001],
          bounds=Bounds([0.0,0.0,0.0,0.0,-np.inf],[1.0, 1.0, np.inf,np.in
f,np.inf]))
```

```
Out[201]:      fun: 3.7367261254672646
      hess_inv: <5x5 LbfgsInvHessProduct with dtype=float64>
      jac: array([ 12.00000002,   4.06740279, 237.14949782,  41.8678734
8,
      -57.39448903])
      message: b'ABNORMAL_TERMINATION_IN_LNSRCH'
      nfev: 432
      nit: 11
      njev: 72
      status: 2
      success: False
      x: array([ 4.99402509e-01,  7.99328776e-01,  0.00000000e+00,
3.93384334e-03,
      -6.09039431e-04])
```

In [202]: *# Now add the minimization to the class*

```
class CurveLTA(object):

    _par_names = ['m0', 'lam', 'betaE', 'betaA', 'betaT']
    _default_pars = [0.5, 0.8, 0.01, 0.005, 0.0001]
    _lower_bounds = [0.0, 0.0, 0.0, 0.0, -np.inf]
    _upper_bounds = [1.0, 1.0, np.inf, np.inf, np.inf]

    def __init__(self):
        pass


    # intializes with some default parameters
    def initialize(self):
        self.m0, self.lam, self.betaE, self.betaA, self.betaT = self._default_pars

    def fit(self, actual, timestamps, mood):
        def loss_func(par):
            self.m0, self.lam, self.betaE, self.betaA, self.betaT = par
            mood_pred = self.predict(actual, time)
            return np.nansum(np.abs(mood - mood_pred))
        res = minimize(loss_func, self._default_pars,
                       bounds=Bounds(self._lower_bounds, self._upper_bounds))
        self.m0, self.lam, self.betaE, self.betaA, self.betaT = res.x
        return res

    def predict(self, actual, timestamps):
        n_trials = len(actual)
        # holds the predicted moods
        mood_pred = np.zeros(n_trials)
        # Holds the exponentially weighted sums for E(t) and A(t)
        sum_E = 0
        sum_A = 0
        for trial_no in range(n_trials):
            if trial_no == 0:
                lte = 0
            else:
                lte = np.mean(actual[:trial_no])
                sum_E = sum_E * self.lam + lte
                sum_A = sum_A * self.lam + actual[trial_no]
                mood_mu = self.m0 + self.betaE * sum_E + self.betaA * sum_A
                + self.betaT * timestamps[trial_no]
                mood_pred[trial_no] = mood_mu
        return mood_pred
```

```
In [203]: CL = CurveLTA()  
CL.fit(actual, time, mood)  
plt.plot(time, CL.predict(actual, time), c="r")  
plt.scatter(time, mood, c="b")
```

```
Out[203]: <matplotlib.collections.PathCollection at 0x7fc0e0141820>
```



```
In [204]: # MAE  
np.nanmean(np.abs(mood - CL.predict(actual, time)))
```

```
Out[204]: 0.10990370957256661
```

## Regularized curve-fitting

Now we add L1 penalties to the coefficients

In [205]: *# Now add the minimization to the class*

```

class CurveLTA(object):

    _par_names = ['m0', 'lam', 'betaE', 'betaA', 'betaT']
    _default_pars = [0.5, 0.8, 0.01, 0.005, 0.0001]
    _lower_bounds = [0.0, 0.0, 0.0, 0.0, -np.inf]
    _upper_bounds = [1.0, 1.0, np.inf, np.inf, np.inf]
    pen_betaE = 0.0
    pen_betaA = 0.0
    pen_betaT = 0.0

    def __init__(self, pen_betaE = 0.0, pen_betaA = 0.0, pen_betaT = 0.0
):
        self.pen_betaE = pen_betaE
        self.pen_betaA = pen_betaA
        self.pen_betaT = pen_betaT

        # intializes with some default parameters
        def initialize(self):
            self.m0, self.lam, self.betaE, self.betaA, self.betaT = self._de
fault_pars


        def fit(self, actual, timestamps, mood):
            def loss_func(par):
                self.m0, self.lam, self.betaE, self.betaA, self.betaT = par
                mood_pred = self.predict(actual, time)
                pen_term = np.abs(self.pen_betaE * self.betaE) + \
                    np.abs(self.pen_betaA * self.betaA) + \
                    np.abs(self.pen_betaT * self.betaT)
                return np.nansum(np.abs(mood - mood_pred)) + pen_term
            res = minimize(loss_func, self._default_pars,
                bounds=Bounds(self._lower_bounds, self._upper_bounds))
            self.m0, self.lam, self.betaE, self.betaA, self.betaT = res.x
            return res

        def predict(self, actual, timestamps):
            n_trials = len(actual)
            # holds the predicted moods
            mood_pred = np.zeros(n_trials)
            # Holds the exponentially weighted sums for E(t) and A(t)
            sum_E = 0
            sum_A = 0
            for trial_no in range(n_trials):
                if trial_no == 0:
                    lte = 0
                else:
                    lte = np.mean(actual[:trial_no])
                    sum_E = sum_E * self.lam + lte
                    sum_A = sum_A * self.lam + actual[trial_no]
                    mood_mu = self.m0 + self.betaE * sum_E + self.betaA * sum_A
+ self.betaT * timestamps[trial_no]
                    mood_pred[trial_no] = mood_mu
            return mood_pred

```

```
In [206]: CL = CurveLTA(pen_betaE = 0.1, pen_betaA = 0.1)
CL.fit(actual, time, mood)
plt.plot(time, CL.predict(actual, time), c="r")
plt.scatter(time, mood, c="b")
```

```
Out[206]: <matplotlib.collections.PathCollection at 0x7fc0e0e47c10>
```



```
In [207]: # MAE
np.nanmean(np.abs(mood - CL.predict(actual, time)))
```

```
Out[207]: 0.061526095054160186
```

## Held-out prediction

We fit to the first 40 trials and predict on the remaining


```
In [297]: CL.fit(actual[:40], time[:40], mood[:40])
```

```
Out[297]: fun: 1.690281251212084
hess_inv: <5x5 LbfgsInvHessProduct with dtype=float64>
jac: array([ 4.          , -0.66621279,  1.43200838, 49.4497209
2,
-42.68056697])
message: b'ABNORMAL_TERMINATION_IN_LNSRCH'
nfev: 900
nit: 21
njev: 150
status: 2
success: False
x: array([ 0.49846875,  0.80032608,  0.00993689,  0.
0.00986397])
```



```
In [298]: plt.plot(time, CL.predict(actual, time), c="r")  
plt.axvline(40)  
plt.scatter(time, mood, c="b")
```

Out[298]: <matplotlib.collections.PathCollection at 0x7fc0e310bd90>



```
In [210]: # MAE (training, test)  
errs = np.abs(mood - CL.predict(actual, time))  
np.nanmean(errs[:40]), np.nanmean(errs[40:])
```

Out[210]: (0.04661197359621433, 0.0747830919056676)

## Ila. Likelihood model

Now we add a utility function and jointly model mood with decision-making.

### Data variables:

- $H(t)$  high gamble
- $L(t)$  low gamble
- $C(t)$  deterministic amount
- $A(t)$  actual outcome
- $G(t)$  whether gambling occurred
- $T(t)$  time elapsed

Define  $W(t) = 1$  if  $A(t) = H(t)$  and  $G(t) = 1$ , and  $W(t) = 0$  otherwise.

Define the subjective win probability as

$$p(t) = \frac{\sum_{i=1}^{t-1} W(i)}{\sum_{i=1}^{t-1} G(i)}$$

with  $p(t) = 0.5$  when it would otherwise be undefined.

### Model parameters:

- $M_0 \in (-\infty, \infty)$  baseline logit-mood
- $\lambda \in [0, 1]$  discount factor
- $\beta_E \in [0, \infty)$  coefficient for LTA
- $\beta_A \in [0, \infty)$  coefficient for actual outcome
- $\beta_T \in (-\infty, \infty)$  coefficient for time trend
- $\gamma \in [0, 2]$  utility function exponent
- $\rho_C \in [0, \infty)$  choice inverse temperature
- $\sigma \in [0, \infty)$  Gaussian noise standard deviation for logit-mood

### Latent variables:

- *Choice bias:*  $V(t) = \rho_C(p(t)H(t)^\gamma + (1 - p(t))L(t)^\gamma - C(t)^\gamma)$
- *Long-term average:*  $E(t) = \frac{1}{t-1} \sum_{u=1}^{t-1} A(u)^\gamma$
- *Predicted logit-mood:*  $\mu(t) = M_0 + \beta_E \sum_{u=1}^t \lambda^{t-u} E(u) + \beta_A \sum_{u=1}^t \lambda^{t-u} A(u)^\gamma + \beta_T T(t)$

### Model:

- $G(t) \sim \text{Binomial}(\text{expit}(V(t)))$
- $M(t) \sim \text{LogitNormal}(\mu(t), \sigma)$

where  $\text{expit}(x) = \frac{1}{1+e^{-x}}$

```
In [211]: from scipy.stats import norm, bernoulli
          from scipy.special import logit, expit
```

```

In [212]: def signedPower(x, y):
            return np.power(np.abs(x), y) * np.sign(x)

class MoodChoiceLTA(object):

    _model_name = 'LTA nonlinear with simple win prob. and choice'
    _par_names = ['m0', 'lam', 'betaE', 'betaA', 'betaT', 'gamma', 'rhoC', 'sigma']
    _default_pars = [0.0, 0.8, 0.01, 0.005, 0.0001, 1.0, 1.0, 0.5]
    _lower_bounds = [-np.inf, 0.0, 0.0, 0.0, -np.inf, 0.0, 0.0, 0.0]
    _upper_bounds = [np.inf, 1.0, np.inf, np.inf, np.inf, 2.0, np.inf, np.inf]

    def __init__(self):
        pass

    # prints parameters
    def __str__(self):
        s = self._model_name
        for par in self._par_names:
            s = s + '\n' + par + ': %.4f' % self.__dict__[par]
        return s

    # initializes with some default parameters
    def initialize(self, params = None):
        if params is None:
            params = self._default_pars
        self.m0, self.lam, self.betaE, self.betaA, self.betaT, \
            self.gamma, self.rhoC, self.sigma = params

    def fit(self, actual, certain, highGamble, lowGamble, choice, timestamps, mood):
        def loss_func(par):
            self.m0, self.lam, self.betaE, self.betaA, self.betaT, \
                self.gamma, self.rhoC, self.sigma = par
            return -self.loglike(actual, certain, highGamble, lowGamble, choice, timestamps, mood)
        res = minimize(loss_func, self._default_pars,
            bounds=Bounds(self._lower_bounds, self._upper_bounds))
        self.m0, self.lam, self.betaE, self.betaA, self.betaT, \
            self.gamma, self.rhoC, self.sigma = res.x
        return res

    def loglike(self, actual, certain, highGamble, lowGamble, choice, timestamps, mood):
        mood_logit, choice_logit = self.predict(actual, certain, highGamble, lowGamble, choice, timestamps)
        choice_ll = bernoulli.logpmf(choice, expit(choice_logit))
        mood_ll = norm.logpdf(logit(mood), loc=mood_logit, scale=self.sigma)
        return np.nansum(choice_ll) + np.nansum(mood_ll)

    def predict(self, actual, certain, highGamble, lowGamble, choice, timestamps):
        n_trials = len(actual)
        # compute the win probabilities

```

```

winIndicator = (highGamble == actual) * choice
pwin = 0.5 * np.ones(n_trials)
temp1 = np.cumsum(winIndicator, axis = 0)
temp2 = np.cumsum(choice, axis = 0)
pwin[temp2 > 0] = temp1[temp2 > 0]/temp2[temp2 > 0]
pwin = np.concatenate([0.5], pwin[:-1])
# holds the predicted moods and choices
mood_logit = np.zeros(n_trials)
choice_logit = np.zeros(n_trials)
# Holds the exponentially weighted sums for E(t) and A(t)
sum_E = 0
sum_A = 0
for trial_no in range(n_trials):
    if trial_no == 0:
        lte = 0
    else:
        lte = np.mean(signedPower(actual[:trial_no], self.gamma
))
        sum_E = sum_E * self.lam + lte
        sum_A = sum_A * self.lam + signedPower(actual[trial_no], self
f.gamma)
        choice_bias = self.rhoC * \
            (pwin[trial_no] * signedPower(highGamble[trial_no], self
.gamma) + \
            (1-pwin[trial_no]) * signedPower(lowGamble[trial_no], s
elf.gamma) - \
            signedPower(certain[trial_no], self.gamma))
        mood_mu = self.m0 + self.betaE * sum_E + self.betaA * sum_A
+ self.betaT * timestamps[trial_no]
        mood_logit[trial_no] = mood_mu
        choice_logit[trial_no] = choice_bias
    return mood_logit, choice_logit

def sample(self, actual, certain, highGamble, lowGamble, choice, tim
estamps):
    mood_logit, choice_logit = self.predict(actual, certain, highGam
ble, lowGamble, choice, timestamps)
    mood_sample = expit(norm.rvs(loc=mood_logit, scale=self.sigma))
    choice_sample = bernoulli.rvs(expit(choice_logit))
    return mood_sample, choice_sample

```

```
In [217]: # use a single subject for all demonstrations
subject_index = 7
time = all_trial_nos[:, subject_index]
actual = all_outcomeAmount[:, subject_index]
mood = all_mood_rating[:, subject_index]
highGamble = all_winAmount[:, subject_index]
lowGamble = all_loseAmount[:, subject_index]
certain = all_certainAmount[:, subject_index]
choice = all_choice[:, subject_index]

MCL = MoodChoiceLTA()
MCL.initialize()
#MCL.predict(actual, certain, highGamble, lowGamble, choice, time)
MCL.loglike(actual, certain, highGamble, lowGamble, choice, time, mood)
```

```
Out[217]: -99.82923444303935
```

```
In [218]: MCL.fit(actual, certain, highGamble, lowGamble, choice, time, mood)
```

```
Out[218]:      fun: 40.08111367796036
      hess_inv: <8x8 LbfgsInvHessProduct with dtype=float64>
      jac: array([ 0.00021743,  0.0007482 , -0.00091873,  0.00686526, -
0.01013589,
      0.00192344,  0.00022311,  0.00174936])
      message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'
      nfev: 981
      nit: 93
      njev: 109
      status: 0
      success: True
      x: array([-1.97364268,  0.2770499 ,  0.4732268 ,  0.11227632,
0.02722128,
      0.56262762,  2.21865169,  0.37428349])
```

```
In [219]: MCL.loglike(actual, certain, highGamble, lowGamble, choice, time, mood)
```

```
Out[219]: -40.08111367796036
```

```
In [220]: print(MCL)
```


```
LTA nonlinear with simple win prob. and choice
m0: -1.9736
lam: 0.2770
betaE: 0.4732
betaA: 0.1123
betaT: 0.0272
gamma: 0.5626
rhoC: 2.2187
sigma: 0.3743
```

## Parameter recovery example

```
In [221]: # Simulate data
mood_s, choice_s = MCL.sample(actual, certain, highGamble, lowGamble, choice, time)
# copy missing pattern of original
mood_s[np.isnan(mood)] = np.nan
```


```
In [222]: # plot original vs simulated mood
plt.plot(time[~np.isnan(mood)], mood[~np.isnan(mood)])
plt.plot(time[~np.isnan(mood)], mood_s[~np.isnan(mood)])
```

Out[222]: [



```
In [223]: # plot choice
plt.plot(time, choice)
plt.plot(time, choice_s)
```

Out[223]: [



```
In [224]: # fit on simulated
MCL2 = MoodChoiceLTA()
MCL2.fit(actual, certain, highGamble, lowGamble, choice_s, time, mood_s)
```

```
Out[224]:      fun: 46.168111630687804
      hess_inv: <8x8 LbfgsInvHessProduct with dtype=float64>
      jac: array([-0.00100471, -0.00409059, -0.00233698, -0.00676579,
0.00619877,
      -0.00301412,  0.00015561, -0.00622364])
      message: b'CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH'
      nfev: 1305
      nit: 124
      njev: 145
      status: 0
      success: True
      x: array([-2.17761285,  0.10234372,  1.09588124,  0.14728704,
0.02959408,
      0.37319938,  2.62004156,  0.41533034])
```

```
In [225]: # simulated
print(MCL2)
```

```
LTA nonlinear with simple win prob. and choice
m0: -2.1776
lam: 0.1023
betaE: 1.0959
betaA: 0.1473
betaT: 0.0296
gamma: 0.3732
rhoC: 2.6200
sigma: 0.4153
```

```
In [226]: # original
print(MCL)
```

```
LTA nonlinear with simple win prob. and choice
m0: -1.9736
lam: 0.2770
betaE: 0.4732
betaA: 0.1123
betaT: 0.0272
gamma: 0.5626
rhoC: 2.2187
sigma: 0.3743
```

## IIb: Bayesian model

We now add prior distributions to the previous model parameters

```
In [227]: from RunDEMC import Model, Param, dists, calc_bplic, joint_plot
```

```

In [241]: # Priors for parameters
params = [Param(name='m0',
                 display_name=r'm0',
                 prior=dists.normal(0, 1)),
          Param(name='lam',
                 display_name=r'$\lambda$',
                 prior=dists.normal(0, 1.4),
                 transform=dists.invlogit
                 ),
          Param(name='betaE',
                 display_name=r'beta_E',
                 prior=dists.normal(-3, 0.7),
                 transform=np.exp,
                 inv_transform=np.log),
          Param(name='betaA',
                 display_name=r'beta_A',
                 prior=dists.normal(-3, 0.7),
                 transform=np.exp,
                 inv_transform=np.log),
          Param(name='betaT',
                 display_name=r'beta_T',
                 prior=dists.normal(0, 0.00001)),
          Param(name='gamma',
                 display_name=r'$\gamma$',
                 prior=dists.gamma(1.5, 0.5),
                 ),
          Param(name='rhoC',
                 display_name=r'$\rho_C$',
                 prior=dists.gamma(1.5, 0.5),
                 ),
          Param(name='sigma',
                 display_name=r'sigma',
                 prior=dists.exp(1))
          ]

```

```

In [242]: # use a single subject for all demonstrations
subject_index = 7
time = all_trial_nos[:, subject_index]
actual = all_outcomeAmount[:, subject_index]
mood = all_mood_rating[:, subject_index]
highGamble = all_winAmount[:, subject_index]
lowGamble = all_loseAmount[:, subject_index]
certain = all_certainAmount[:, subject_index]
choice = all_choice[:, subject_index]

```



```
In [243]: def eval_fun(params):
            md = MoodChoiceLTA()
            params2 = np.array([params[n] for n in md._par_names])
            n_params, n_particles = params2.shape
            ll = -np.inf * np.ones(n_particles)
            valid1 = (params2 > np.reshape(np.array(md._lower_bounds), (n_params
, 1)))
            valid2 = (params2 < np.reshape(np.array(md._upper_bounds), (n_params
, 1)))
            valid = np.logical_and(valid1, valid2)
            for ind_part in np.nonzero(valid)[0]:
                md.initialize(params2[:, ind_part])
                ll[ind_part] = md.loglike(actual, certain, highGamble, lowGamble
, choice, time, mood)
            return ll
```

```
In [244]: m = Model('mmi', params=params,
                    like_fun=eval_fun,
                    init_multiplier = 3,
                    use_priors = False,
                    verbose=True)
```


```
In [245]: m._initialize(num_chains=50)
```

```
Initializing: 150(50) 145(45) 142(42) 139(39) 136(36) 132(32) 126(26) 1
20(20) 117(17) 113(13) 110(10) 106(6) 103(3)
```

```
In [246]: np.min(m.weights[-1]), np.max(m.weights[-1])
```

```
Out[246]: (-17322870.812327046, -68.18925915511545)
```

```
In [247]: plt.figure(figsize=(12,10))
for i in range(8):
    plt.subplot(3,4,i+1)
    plt.hist(m.particles[:, :, i].flatten(), bins='auto', alpha=.5, density=True)
    plt.title(MCL._par_names[i])
```




```
In [248]: times = m.sample(100, burnin=True, migration_prob = 0.1)
```

```
Iterations (100): 1 2 3 4 5 6 7 8 x 9 10 11 12 13 14 15 16 17 18 19 20
21 22 23 24 25 26 27 28 29 30 31 32 33 34 x 35 36 37 38 39 40 41 42 43
44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67
68 69 70 71 72 73 x 74 x 75 76 77 x 78 79 80 81 82 83 84 85 86 x 87 88
89 90 91 92 93 94 95 96 97 98 99 100
```

```
In [249]: np.min(m.weights[-1]), np.max(m.weights[-1])
```

```
Out[249]: (-101.60792264263635, -68.18925915511545)
```


```
In [250]: plt.figure(figsize=(12,10))
for i in range(8):
    plt.subplot(3,4,i+1)
    plt.hist(m.particles[0, :, i].flatten(), bins='auto', alpha=.5, density=True)
    plt.hist(m.particles[-1, :, i].flatten(), bins='auto', alpha=.5, density=True)
    plt.title(MCL._par_names[i])
```




```
In [251]: times = m.sample(100, burnin=False)
```

```
Iterations (100): 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45
46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69
70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93
94 95 96 97 98 99 100
```

```
In [252]: # show the chains are mixing and converging
plt.plot(m.weights[-100:], alpha=.3);
```



```
In [253]: plt.figure(figsize=(12,10))
for i in range(8):
    plt.subplot(3,4,i+1)
    plt.hist(m.particles[0, :, i].flatten(), bins='auto', alpha=.5, density=True)
    plt.hist(m.particles[-1, :, i].flatten(), bins='auto', alpha=.5, density=True)
    plt.title(MCL._par_names[i])
```



```
In [261]: # debugging
```

```
In [289]: p1=m.particles[-1, 0]
p1
```


```
Out[289]: array([-9.97538338e-05,  6.00347560e-01,  5.43897946e-02,  5.24324942e-02,
                1.91922978e-05,  5.45469232e-01,  2.89763620e+00,  7.88518713e-01])
```

```
In [290]: p2=m.particles[-1, 3]
p2
```

```
Out[290]: array([-7.31339565e-01,  5.44584541e-01,  1.55575456e-02,  7.67690370e-02,
                4.10651593e-05,  6.73260648e-01,  2.64165594e+00,  1.15006514e+00])
```

```
In [291]: plt.plot(m.weights[-100:, [0,3]])
```

```
Out[291]: [<matplotlib.lines.Line2D at 0x7fc0e0ec0eb0>,  
<matplotlib.lines.Line2D at 0x7fc0e0ec0fa0>]
```



```
In [292]: def ll(params):  
            MCL = MoodChoiceLTA()  
            MCL.initialize(params)  
            return MCL.loglike(actual, certain, highGamble, lowGamble, choice, t  
ime, mood)
```

```
In [296]: [ll(p) for p in m.particles[-1]]
```

```
Out[296]: [-68.55036214985282,  
-80.854936970193,  
-69.23349712768768,  
-69.29013840688418,  
-78.90207021908176,  
-70.29094947429809,  
-70.0605273538265,  
-69.23349712768768,  
-80.81496092937205,  
-69.23349712768768,  
-78.90207021908176,  
-78.90207021908176,  
-78.90207021908176,  
-69.23349712768768,  
-69.23349712768768,  
-69.23349712768768,  
-69.23349712768768,  
-69.23349712768768,  
-78.84060454146022,  
-78.90207021908176,  
-69.23349712768768,  
-100.71895772569297,  
-78.29758245238412,  
-69.23349712768768,  
-78.29758245238412,  
-69.23349712768768,  
-78.90207021908176,  
-69.23349712768768,  
-78.90207021908176,  
-69.23349712768768,  
-80.81496092937205,  
-80.81496092937205,  
-69.23349712768768,  
-78.90207021908176,  
-80.81496092937205,  
-78.90207021908176,  
-78.90207021908176,  
-69.23349712768768,  
-78.90207021908176,  
-80.81496092937205,  
-78.90207021908176,  
-78.84060454146022,  
-69.23349712768768,  
-69.23349712768768,  
-78.90207021908176,  
-78.90207021908176,  
-78.90207021908176,  
-69.23349712768768,  
-78.90207021908176,  
-69.23349712768768]
```

```
In [293]: ll(p1)
```

```
Out[293]: -68.55036214985282
```

In [294]: `ll(p2)`

Out[294]: `-69.29013840688418`

In [295]: `ll(0.5 * p1 + 0.5 * p2)`

Out[295]: `-64.68052793018006`

In [ ]: